

# STA304 - Fall 2022

## Assignment 1

Eric Zhu - 1005131368

30/09/2022

### Part 1

#### Goal

A Graphics Processing Unit is a microprocessor that is suited for certain kinds of computations. Fundamentally, a GPU is a SIMD device (simultaneous instruction multiple data) device that enables much faster computation of matrix heavy operations, e.g., computer graphics and certain kinds of numerical optimization algorithms (gradient descent). They were originally built to accelerate rasterization performance (image rendering) across computer rendering with GPU cores that performed much faster than CPU cores for rendering thanks to task specialization [3]. Regardless of their origins, traditional GPU cores have been fundamental to machine learning experimentation, scientific modelling, and entertainment.

In recent years, GPUs have gotten more capabilities via specialized hardware for tasks like real-time ray tracing (simulating realistic lighting) and half-precision (16 bit) matrix math [1]. Consumers have found this to be of dubious value depending on their background and use cases, leading to discussions across the industry and community about what functions GPUs should prioritize [2]. As the GPU market continues its strong growth, slated to become a 200.85 billion dollar market by 2027 [4], the question over GPU functionality preferences also grows in importance and relevance. While there may be many preferences for GPU capabilities, the preferences we wish to focus on are based on the distinct hardware capabilities that are most relevant across consumer categories, i.e., traditional GPU cores for rasterization/compute capabilities, real-time ray-tracing hardware, and hardware for half-precision matrix math (will also be referred to as half precision hardware). Essentially, we are aiming to examine the question: what kind of GPU user wants what kind of GPU hardware?

#### Procedure

Fundamentally, our goal is to analyse the preferences of GPUs across the consumer categories, including a diverse set of users with wildly different use cases. Thus, our target population is simply the users of GPUs across these different categories, and specifically for this survey, we are looking at two aggregated categories (professionals and consumers). The frame population is dependent on the likely ways of distributing this survey. The easiest way of distributing this survey would be through online forums of GPU users. Other ways would have significant drawbacks. For example, distributing it through exit surveys after GPU purchasing would leave us a biased sample from those who use GPUs but do not directly purchase or install them, e.g., academics who run experiments on GPUs. Or the other option of using customer data collected by “game optimization” software from the manufacturer (like GeForce Experience) would bias our sample given that it is generally uncommon for enterprise settings to install these apps (e.g., AWS does not install these) and would potentially lead to biases for certain GPU manufacturers. Since forums are targeted to a specific demographic, it would make the most sense for us to employ stratified random sampling as this strategy

follows most naturally from targeted forums where we would distribute our survey as they naturally form self-selecting strata. This is advantageous over other sampling methods, namely cluster random sampling. This is because it is simply unfeasible to form representative clusters with online forums, i.e., forum membership is voluntary and forum topic is predetermined. In other words, there is no good way of getting lists of users from forums like Reddit where such information is private. Thus the only feasible way would be to scrape users who comment and post as those are the only two publicly available pieces of user information. Then we'd have to send those users messages about our survey (after forming clusters), which could introduce bias should we use them as our sample since users who post generally are more vocal or passionate about some subject. So cluster random sampling could introduce worrying sources of sample bias. While stratified random sampling is definitely the better option among others, the major downside is that members of certain strata may belong in many strata, e.g., it is reasonable to assume that there are overlaps between gamers and academics (perhaps some PhD student plays video games). This may introduce some sampling bias as we might over/under represent some groups; this is unavoidable even with a different sampling technique as this source of bias is inherent to the individual so this bias would still be present regardless of sampling technique.

### Showcasing the survey.

Survey link: <https://forms.office.com/r/uqb23eNtc7>

I chose the following three questions overall because the first two are representative of questions that would enter the models as covariates, while the third question would enter the model as a response variable.

**Question 1: Estimated hours per day of intensive GPU usage? Intensive usage could be an application like training Deep Learning models, running 3D graphics applications like games, and or rendering videos.**

This question is meant to gather information measuring the participant's familiarity of GPUs. The logic here is that the longer a user uses a GPU for their specified use case (e.g., a deep learning researcher interfacing with GPU based deep learning libraries), the more familiar that user would be with GPUs and their capabilities. This is similar to the logic how we could reasonably conclude that a person is familiar the functionality of their laptop should they use it a lot. An inherent drawback to this question is that it doesn't ask directly about a user's perceived familiarity with GPUs (perhaps on some scale, e.g., 1 to 10), so it's a proxy for it. But this is done because perceived familiarity can be heavily biased by their own perception, so this provides a more absolute measure of familiarity with GPUs. Another potential drawback is that since we allow for any input, this question measures the estimated minutes on a continuous scale; perhaps a discrete scale could provide us with the ability to model this variable using random effects, but this was done because there's no good cutoff for each group. In other words, a cutoff would be misinformed and could possibly lead to skewed analysis.

**Question 2: What type of GPU do you usually use?**

Options:

- Datacentre/HPC (Nvidia A series, AMD instinct)
- High end consumer (Nvidia xx80 (ti), Nvidia xx90(ti), AMD x800, x800xt, x900xt, e.g., Nvidia RTX 3080)
- Midrange consumer (Nvidia xx70 (ti), Nvidia xx60(ti), AMD x700, x700xt, x600xt, e.g., Nvidia RTX 3070)
- Entry level consumer (Otherwise: Nvidia xx50 and below, AMD x500 and below, e.g., Nvidia RTX 3050)

This question is meant to gather information measuring the effect on GPU capability preferences due to a user's hardware segmentation. In other words, we are looking to gather information on if using a certain

kind of GPU would affect their preferences for GPU hardware capabilities because different “levels” of GPUs have different characteristics. The logic here is comparable to asking a car owner about their current car in a survey about car feature preferences. Additionally, a strong point is that the categorization of the options is based on industry standards. As in, the datacentre/HPC categorization and corresponding examples (Nvidia A series, AMD instinct) are based on how they are marketed by both companies. For the consumer options (last 3), those are based on search results/a listing from Best Buy, a major electronics retailer [12] [13]. This should mean that there is very little ambiguity (given our exhaustive options listing) for the survey taker in choosing their appropriate category since this categorization would be exactly the same categories GPU buyers are exposed to. The final major pro of this question is that it is implicitly ordered, i.e., the categories are ordered by price, so we could examine how different GPU price categories affect preferences too. However, the downside to this is that some users may not necessarily agree with these categories, so they may not choose to follow the examples in the parenthesis as stated in the four options.

**Question 3: How important are traditional GPU cores in GPUs? Consider 50 to be neutral and 100 to be a strong like; negative scores will indicate a strong dislike. Decimals are allowed.**

This question is meant to gather information on a user’s preference for traditional GPU cores, which is the response variable for one of the models. This question is repeated three times for the three hardware categories, i.e., traditional GPU cores, ray-tracing hardware, and half-precision matrix math (half-precision) hardware. We allow for decimal values so that it is appropriate later to model the response with a continuous probability distribution. So this question is meant to be interpreted along side the other two, and this is a major strength over making the user choose a particular specialized hardware preference or even a ranking. Primarily, this will allow us to examine all three hardware categories independently should the same covariates affect the response variable differently for each hardware category, i.e., we can get separate effect size estimations for the same covariates but for different response variables. A drawback is that this scale is not rigorously defined, and is the most open ended question in the survey. This is because GPU users may not consider the same scale as the other users, e.g., 50 may be neutral to some users but may indicate a slight preference to others. To remedy this I added the “and 50 being neutral” part of the question. Nonetheless, users will likely still consider the scale to be slightly different based on their own experiences. A discrete scale, e.g., categories of preferences, could have been better as it would be more rigorously defined (labeled categories), but there would still be some bias regardless since the source of bias is the individual’s conception of the scale, i.e., how they conceptualize favour for some hardware capability.

## Part 2

### Data

We are looking to simulate the data due to practical sampling constraints. We have a few main considerations: the demographics that we are looking to simulate, and the effects stemming from the covariates we wish to measure. At a high level, we are looking to measure the preferences of individuals, given their use cases, for the three kinds of hardware in GPUs we are considering: GPU cores, ray-tracing hardware, and half-precision matrix math hardware. Since it is conceivable that certain variables from our survey, e.g., type of GPU user will affect the preferences overall for a user type, we will consider the data generation in a so-called hierarchical manner.

First, we need to simulate our target population, i.e., what are the proportions of our GPU user types? While the number of each GPU user type wouldn't be released as public information (think military users or sensitive corporate situations like Amazon AWS), we can inform this simulation choice using publicly available earnings data from the two major GPU makers: Nvidia and AMD. In both cases, their public earnings data show that they sell roughly equally to the consumer and professional markets in Q1 of 2022 [5][6]. Thus, we will first sample from the Bernoulli distribution with a parameter  $p = 0.5$ , i.e., it is equally likely for us to choose a professional or a consumer GPU user.

Second, with a user category chosen, we then need to simulate the next three survey questions that will act as covariates in our model, i.e., factors that could affect GPU hardware category preferences. Recall that the last 3 questions are essentially response variables that ask users for their preferences of GPU hardware categories.

The second question in our survey asks the user their primary interaction with their GPU. The question accepts three options: programming libraries (like Pytorch), high level graphical interfaces (like Photoshop), and graphics applications (like 3D games). The options are implicitly ordered by “involvedness” of the user's primary interaction with their GPU; in other words, we can provide an ordering to the options based on how abstracted the user's primary interaction is from dealing with the GPU. Since the ordering is based on how hands on the user's primary interaction is, we have the ordering be that programming libraries are the most hands on (a label of 3), followed by graphics applications (a label of 2), then high level graphical interfaces (a label of 1). In other words, programming libraries are least abstracted and high level graphical interfaces are most abstracted away from the GPU. Since we do not have access to how people use their GPUs, we will have to come up with some sensible values for the parameter of a categorical distribution to simulate this question. Recall that a categorical distribution takes a vector of probabilities, which in our case is dependent on the GPU user type (either professional or consumer) and is of length 3. If our user is a professional, it is sensible that the 60% of professionals primarily interact with GPUs through programming libraries (ML researchers, academics, software engineers, etc...), then 30% of professionals interact with GPUs through high level interfaces (creatives who use apps like Photoshop), and then 10% of professionals interact with GPUs through graphics applications (video game designers and professionals who play a lot of video games). Then for consumers, it is sensible that 80% of consumers interact with GPUs through graphics applications (primarily video games as GPUs were initially made for them), then 10% primarily use them for programming libraries (those who may use them for ML personal projects), and then 10% primarily interact with them through high level graphical interfaces (those who have them for editing photos or videos). So to recap, we will simulate this variable by sampling from a `Categorical([0.6, 0.3, 0.1])` distribution if they are a professional and a `Categorical([0.1, 0.1, 0.8])` distribution if they are a consumer.

The third question asks for the type of GPU the user uses. We unfortunately do not have data from professionals about the distribution of their GPU types since in many cases they may be custom or trade secret information. However, from having used the University of Toronto Computer Science department GPU clusters as part of projects in the past, an institution like the University of Toronto uses almost exclusively high end or datacentre class GPUs. The CS department website also provides an example list of GPUs that are part of a general GPU cluster at UofT [8]. With that in mind, we will simulate this variable given a professional GPU user using the categorical distribution where 50% of professionals use datacentre GPUs, 40% use high end consumer models, 8% use midrange consumer models, and 2% use low end consumer models. So

provided a professional GPU user we sample from `Categorical([0.5, 0.4, 0.08, 0.02])`. Alternatively, provided a consumer GPU user, we can look at the Steam hardware survey to gather some insights. Steam is a major platform used for buying/selling video games, accounting for 3.1 billion dollars in sales for the first half of 2022 [9]. While these estimates are biased towards gaming, it is the best large scale estimate of consumer GPU types. Additionally, information is collected if the Steam user opts in; it is not indicative of if the Steam user is primarily a gamer. From the Steam hardware survey, we gather that essentially no consumer uses a datacentre/HPC GPU, high end consumer GPUs account for approximately 10% of consumer GPUs, and the rest is roughly evenly split between midrange and entry level consumer GPUs. So given a consumer, we will sample this variable from `Categorical([0, 0.1, 0.45, 0.45])`.

Then the fourth question asks about the estimated hours per week of GPU usage. This variable, as mentioned previously, is a proxy for a user’s familiarity with GPUs as the longer they use a GPU through GPU based tasks (like gaming) the more familiar they should be with GPUs. This is also hard to simulate since a user’s GPU usage is private information. However, we can infer that it would be sensible that professional users of GPUs use them for approximately 40 hours per week. So then, if we were to generate samples for this variable, we could use a normal distribution with a mean of 40 and a standard deviation of 5 (to account for overtime or leaving early each work day). Then for consumers, we can look to the “The State of Online Gaming – 2021 survey”, which surveyed 4000 people in 8 countries. The survey was conducted by Limelight, a cloud service provider. On average, they found that an average gamer played 8 hours and 27 minutes of video games per week [10]. Additionally, the top 25% of gamers play more than 12 hours per week. Using 12 hours per week as the 75<sup>th</sup> percentile (0.674 standard deviations from the mean), we can use the sample standard deviation as the estimate for the population standard deviation, which is  $\approx 5.2$  hours when solving for  $0.674 \cdot \hat{\sigma}_{sample} = 3.5$  [10]. So then, if we were to generate samples for this variable, we could use a normal distribution with a mean of 9 and a standard deviation of 5.2. To recap, given a professional user, we will sample from  $\mathcal{N}(40, 5)$ , and provided a consumer, we will sample from  $\mathcal{N}(9, 5.2)$ . Finally, we will absolute all of the samples should there be a negative sample as negative hours is nonsensical.

Finally, we will need to weight these simulated variables to compute a conditional expectation for the three response variables. Then, we will use this mean as the mean to a normal distribution. We will also use a constant variance parameter (one  $\sigma^2$  for all  $x_i$ ’s) because we do not have data to allow us to faithfully make simulation choices for different  $\sigma^2$ .

As mentioned previously, we should model the GPU user type hierarchically, meaning that we find it reasonable that the type of user affects the baseline for GPU hardware category preference scores. We will set both the professional and consumer baseline preference for traditional GPU cores as 60 because they are used in almost all GPU applications [3]. Then we will set the baseline preferences for ray-tracing hardware as much higher in consumers than professionals as that hardware is really only used in video games, so we will set the baseline preference of consumers to 50 (neutral) and 10 for professionals [2]. Finally, we set the baseline preferences for half-precision matrix math hardware as both 55 for consumers and professionals because they are used in AI task acceleration, which is rapidly used more and more in both consumer applications (video game upscaling [18]) and professional applications (deep learning training [19]).

We will first consider the weights for the second survey question (a user’s primary interaction with their GPU). As mentioned previously, the primary interaction questions are ranked in order from least interaction to most interaction. Thus, we need to choose a weight for this question that is sensical to interpret with respect to preference scores and other levels. We choose 4 to be our weight because since we have 3 levels, the “top” level, i.e., programming libraries, would have a total contribution of 16 to the total preference score, which is reasonable. This is because the way a user interacts with their GPU should be pretty informative for how perceive it. Additionally, we add a negative constant of -2 to the baseline weight of 4 for consumer users when it comes to ray tracing and half-precision hardware because those pieces of hardware are more advanced, and consumer users, on average, likely use them less.

Then we consider the weights for the third survey question (type of GPU). Also recall that the options are ordered in terms of price for these GPUs, so like the previous we need to choose a weight that is sensical for the 4 levels we have, i.e, our top level will have a 4 times contribution to the final preference score over the baseline level. We choose the weights to be 2 for the traditional GPU cores and half-precision matrix math hardware categories, but 0.5 for the ray tracing category. This is representative of the fact that ray-tracing

isn't nearly as important as traditional hardware currently. For ray tracing hardware, we choose 0.5 since real-time ray tracing is still rather niche [2]. We once again adjust these weights for consumers because they have a different distribution of GPUs, owing to their use case. So we choose weights for the traditional GPU cores and half-precision matrix math hardware categories to be 1 each. This is representative of consumer use cases, e.g., video games, where, across the board, experiences are more curated relative to professionals.

Finally, we consider the weights for the fourth survey question (hours of usage per week). We will weight this effect for all three categories as 0.25 since each individual hour sensibly contributes very little to a user's overall perception of GPU capabilities. Additionally, we find no reason to suggest that a professional user's hour GPU use would be weighted differently than a consumer's. Indeed, they use their GPUs differently, but the usage is meaningful to both categories, which is then reflected in their preference scores for different hardware.

Then the individual computed conditional expectations were used as the mean for a normal distribution. Here we use a constant standard deviation for all conditional expectations simply because there is no data to otherwise inform a different modelling choice. We use a standard deviation of 5 because this allows for some amount of variation, which should represent individual variation due to sources like question interpretation/interpretations of scale (how much dislike must a user have to warrant a negative score) despite answering the same on survey questions.

For this simulated dataset, we simulated 1000 conditional means using the seed 123, then for each conditional mean, 10 samples were generated from the parameterized normal distribution. So overall, there are 10000 samples in this simulated dataset.

**Note:** Since the data was simulated, there was no data cleaning involved.

Below is a table of the important variables from our simulated dataset:

Table 1: The table of the important variables of simulated dataset

Variable	Description
<b>user</b>	The binary encoded user type, i.e., professional (0) and consumer (1)
<b>GPUType</b>	Integer encoded type of GPU (ordered by price), i.e., datacentre/HPC (4), high end consumer (3), midrange consumer (2), entry level (1)
<b>primaryInteraction</b>	Integer encoded user's primary interaction with a GPU (ordered by use case's abstraction from GPU) , i.e., programming libraries (3), graphics applications (2), high level graphical
<b>weeklyHours</b>	The number of hours per week that the user uses their GPU
<b>simulations.tradGPUCores</b>	The simulated preference for traditional GPU cores given other simulated survey responses
<b>simulations.rayTraced</b>	The simulated preference for ray tracing hardware given other simulated survey responses
<b>simulations.halfPrecision</b>	The simulated preference for half precision matrix math hardware given other survey response

Below is the five number summary of the simulated traditional GPU cores preference scores:

Table 2: Five number summary of simulated traditional GPU cores preference scores

Min.	1st Quartile	Median	Mean	3rd Quartile	Max
54.46	73.93	77.85	77.86	81.74	100.64

Then, we have also have a five number summary of the simulated ray tracing hardware preference scores:

Table 3: Five number summary of simulated ray tracing hardware preference scores

Min.	1st Quartile	Median	Mean	3rd Quartile	Max
8.379	26.481	39.663	42.828	59.269	78.321

Finally, we have a five number summary of the simulated half precision matrix math hardware scores:

Table 4: Five number summary of simulated half precision hardware preference scores

Min.	1st Quartile	Median	Mean	3rd Quartile	Max
46.91	65.19	69.87	70.20	75.01	95.80

From the three five number summaries, we see that the distribution of the simulated traditional GPU cores preference scores had the least variation as indicated by the IQR, the simulated half precision matrix math hardware preference scores has the second most variation, and simulated ray tracing hardware preference scores had the least. Similarly, simulated traditional GPU cores preference scores had the highest measures of centre (median and mean), simulated ray tracing hardware preference scores had the second most, and simulated half precision matrix math hardware had the least.




Figure 1: Weekly Hours of GPU Usage vs Traditional GPU Cores Preference Score

From the scatter plot of **Weekly Hours of GPU Usage** versus **Traditional GPU Cores Preference Score**, we see that both distributions are approximately Gaussian; however the consumer distribution is not exactly Gaussian as weekly hours of GPU usage must be greater than 0. Additionally, the Professional category of users has a generally higher preference for traditional GPU cores. It appears that the difference between the two user groups is just a shift up/down, i.e., consumers appear to have been shifted down by some constant. The spread between the distributions appear to be similar in general.




Figure 2: Weekly Hours of GPU Usage vs Ray Tracing Hardware Preference Score

From the scatter plot of **Weekly Hours of GPU Usage** versus **Ray Tracing Hardware Preference Score**, we see that both distributions are also approximately Gaussian; however the consumer distribution is not exactly Gaussian as weekly hours of GPU usage must be greater than 0. Additionally, the Consumer category of users has a generally higher preference for ray tracing hardware. It appears that the difference between the two user groups is just a shift up/down, i.e., professionals appear to have been shifted down by some constant. The spread between the distributions appear to be similar in general.




Figure 3: Weekly Hours of GPU Usage vs Half Precision Hardware Preference Score

From the scatter plot of **Weekly Hours of GPU Usage** versus **Half Precision Matrix Math Hardware Preference Score**, we see that both distributions are also approximately Gaussian; however the consumer distribution is not exactly Gaussian as weekly hours of GPU usage must be greater than 0. Additionally, the Professional category of users has a generally higher preference for half precision matrix math hardware. It appears that the difference between the two user groups is just a shift up/down, i.e., consumers appear to have been shifted down by some constant. The spread between the distributions appear to be similar in general.

There are additional plots in the appendix that depict the scatterplots above, but separated by the corresponding GPU type and primary interaction with their GPU. From the additional scatter plots, we can see that there are fairly noticeable changes in the slope of the scatter plot



**Note about simulation:** The data was simulated specifically using the R language with base R packages and the R library `extraDistr` for sampling from a categorical distribution [20].

## Methods

First from our plots, we see that in some cases, there is a pretty clear visual distinction between the preferences of GPU hardware capabilities of professional versus consumers. In particular, ray-tracing hardware seems to have a very clear difference. This indicates that potentially we will need to employ a model that considers this, i.e., a mixed effects model. At a high level, a mixed effects model is just one that allows for more granularity to the model coefficients by allowing for the coefficient of each variable to be adjusted based on some other categorical variable or for the intercept to be adjusted based some other categorical variable. While other flexible models exist, a mixed effects model is much more interpretable (consider a neural network with thousands of parameters versus a mixed effects model with a few). So our modelling approach was a hierarchical one as we believe that the preference scores are affected by an overarching variable, i.e., the user type. To gather evidence of this hierarchical nature, we calculated a 95% confidence interval for the difference in means (using unequal variances) in independent samples. So we made sure it is appropriate to do so, i.e., checked for normality (if the data conforms to the bell shaped normal distribution) using QQ-plots, and we employed the F-test for equal variances.

As a note about investigating normality, we will not use normality tests. This is because the standard Shapiro-Wilk test is too sensitive on large sample sizes, e.g., over 5000 (which is the sample size limit in R and we exceed that). So even if we have small deviations from normality, we will get evidence of it even if the tests at this sample size, e.g, T-tests, are robust to such deviations. [11] The Kolmogorov-Smirnov test is another common one, but that test has low power and we must estimate our distributional parameters from data. Instead QQ-plots are preferred due to easy of interpretation at large sample sizes. [11]

Then for the F-test for equal variances, we used a two tailed F-test using the hypotheses:

$$\begin{aligned} H_0 : \sigma_{\text{Professional GPU User}}^2 &= \sigma_{\text{Consumer GPU User}}^2 \\ H_a : \sigma_{\text{Professional GPU User}}^2 &\neq \sigma_{\text{Consumer GPU User}}^2 \end{aligned}$$

We use the following test statistic:

$$F = \frac{\sigma_{\text{Professional GPU User}}^2}{\sigma_{\text{Consumer GPU User}}^2}$$

Note that the  $F$  test statistic above follows an F distribution with  $n - 1$  and  $m - 1$  degrees of freedom under the null hypothesis. Also since the number of samples between user types are unlikely to be equal, we will randomly sample 1000 samples from each subgroup.

Due to the small p-value then we have evidence that the variances of distributions among the two types of GPU users is not equal, which motivated us to use a confidence interval with unequal variances.

Our confidence interval will compare the difference in means of the distributions of the Professional GPU User scores to those of Consumer GPU User scores. We will do this three times, with one for each category, we are not performing multiple testing and do not need correction (e.g., Bonferroni correction). We opt for a confidence interval over a hypothesis test because we wish to explore the effect size of the difference in the means rather than the point estimate exactly. If the confidence interval (CI) does not contain 0, then we can also conclude that there's a statistically significant difference in the means. We will compute the confidence interval as [14]:

$$\begin{aligned}
(\hat{\mu}_{\text{professional}} - \hat{\mu}_{\text{consumer}}) \pm t^* \sqrt{\frac{s_{\text{professional}}^2}{n_{\text{professional}}} + \frac{s_{\text{consumer}}^2}{n_{\text{consumer}}}} \\
t^* = t_{(\alpha/0.2, df)} \\
df = n_{\text{professional}} + n_{\text{consumer}} - 2
\end{aligned}$$

Our analysis for our results consists of mixed effects models that separately measure the preference for traditional GPU cores, ray-tracing, and half-precision matrix math hardware. In particular, we have three models using the same model specification but different response variables. So we have a model for each one of the preferences for traditional GPU cores, ray-tracing, and half-precision matrix math hardware. We construct these three models separately because we want to be able to examine each response variable, i.e., the preferences of the three categories of hardware separately. This can allow us to see if we get different effect size estimates for the same covariates with a different response variable. Recall that an effect size is simply the coefficient of a model for some variable, which is generally written as  $\beta_k$  for the  $k^{th}$  coefficient/variable. We'll specifically model the response variable using a mixed effects model, where the response follows a normal distribution. So precisely, we are using a linear mixed effects model. We make this modelling choice as the data is continuous and does not appear to be constrained (no need for a distribution like a truncated normal), and from the graphs above, they appear to be approximately normally distributed. Additionally, we add a so-called random effect due to the user type, which, as discussed above, is motivated from the observation that scores are separated by user types. In particular, we add a random intercept, meaning the intercept of the model is adjusted for each user type. We also add random slopes (adjust the coefficients for each variable by the user type) for the **Primary GPU Interaction** and **Type of GPU** variables because as discussed in the survey creation, we believe that different user types would consider the same GPU interaction and GPU type differently, meaning we'd need to adjust the coefficients based on user type. For our primary model we do not add a random slope to the **Estimated Hours of Daily GPU Usage** variable because visually it appears that coefficient for that variable is constant across user types. However we investigate this further with a more complex model and present the results in the results section. We define our primary model as follows:

$$\begin{aligned}
y_i &\sim N(\mu_i, \sigma^2) \\
\mu_i &= \beta_0 + (\beta_1 + U_{1j}) \cdot \text{Primary GPU Interaction} + (\beta_2 + U_{2j}) \cdot \text{Type of GPU}_i \\
&\quad + \beta_3 \cdot \text{Estimated Hours of Daily GPU Usage}_i + U_{3j} \\
U_1 &\sim N(0, \tau_1) \\
U_2 &\sim N(0, \tau_2) \\
U_3 &\sim N(0, \tau_3)
\end{aligned}$$

Note that above,  $\beta_0$  is the intercept, and  $U_{1j}$  is the random slope for the **Primary GPU Interaction** for the  $j^{th}$  user group,  $U_{2j}$  is the random slope for the **Type of GPU** for the  $j^{th}$  user group, and  $U_{3j}$  is the random intercept for the  $j^{th}$  user group. Then  $U_1 \sim N(0, \tau_1)$ ,  $U_2 \sim N(0, \tau_2)$ , and  $U_3 \sim N(0, \tau_3)$ , meaning that these random effects are described by a normal distribution centered around 0 with some standard deviations. The random effects are from distributions from centered around 0 because they may not have an effect for some group, and are given independent standard deviations as this allows for more granularity, e.g., the user type may affect the **Primary GPU Interaction** differently than how it affects **Type of GPU**.

As mentioned above we also investigate if the **Estimated Hours of Daily GPU Usage** variable also needs a random slope for the user type, i.e, if the user type affects **Estimated Hours of Daily GPU Usage** coefficient differently by user type. In this case, we use the same variable but with one additional random slope:

$$\begin{aligned}
y_i &\sim N(\mu_i, \sigma^2) \\
\mu_i &= \beta_0 + (\beta_1 + U_{1j}) \cdot \text{Primary GPU Interaction} + (\beta_2 + U_{2j}) \cdot \text{Type of GPU}_i \\
&\quad + (\beta_3 + U_{4j}) \cdot \text{Estimated Hours of Daily GPU Usage}_i + U_{3j} \\
U_1 &\sim N(0, \tau_1) \\
U_2 &\sim N(0, \tau_2) \\
U_3 &\sim N(0, \tau_3) \\
U_4 &\sim N(0, \tau_4)
\end{aligned}$$

For brevity sake, we will not discuss the model in detail again. The singular addition to this model is  $U_{4j}$ , which is the random slope that enables a different coefficient for **Estimated Hours of Daily GPU Usage** based on user type.  $U_4$  is distributed by a  $N(0, \tau_4)$  distribution for the same reasons above.

Recall that we chose a linear mixed effects model over other flexible choices like neural networks because linear mixed effects models are relatively easy to interpret. However, adding this random slope makes the model more complicated to interpret, so we will use a likelihood-ratio test to determine if it is worth adding this extra complexity. To do so, we will use the R package `lme4` [16].

The likelihood-ratio test examines if the data has a higher likelihood under our more complex model, i.e., if the model is essentially a better fit. If the more complex model is significantly a better fit (has a significantly higher likelihood), then we would choose the more complex model. So the associated hypotheses are:

$$\begin{aligned}
H_0 &: \text{The simpler model is better} \\
H_a &: \text{The more complex model is better}
\end{aligned}$$

So if we have a small p-value from the resulting test, then we have evidence that the more complex model is better. We will opt to not show the derivations of the test statistic as it is rather complex and is out of the scope of this paper. For clarity, the test statistic we use is:

$$-2 \log \left( \frac{\mathcal{L}_{\text{simple model}}(\hat{\theta}_{\text{simple model}} \mid x)}{\mathcal{L}_{\text{complex model}}(\hat{\theta}_{\text{complex model}} \mid x)} \right)$$

Note that  $\mathcal{L}(\hat{\theta} \mid x)$  is the log-likelihood function provided data  $x$  and maximum likelihood estimated parameters  $\hat{\theta}$ .

Finally note that all of these models are fit using the R package `lmer` [15], and model specification as code is available in the results section. All tables for results and in the rest of this paper were prepared using the R package `knitr` [17]. With our final models, we will present the coefficients, and interpret them to analyze the effects of our covariates on our three hardware category preference scores.

## Results

First, we will examine the QQ-plots of all three preference score categories, where each plot contains the QQ-plot for both consumer and professional groups. This as discussed in the methods section will examine if our data has worrying deviations from normality.




Figure 4: QQ-plots for Traditional GPU Cores Preference Scores by User Type

The above plot depicts the QQ-plots for Traditional GPU Cores preference scores. In the consumer subplot, we see that the data mostly conforms to a normal distribution, where there are some deviations in the tails. The deviations in the upper tails are indicative of a very slightly lighter than theoretical upper tail and a very slightly lighter heavier tail. Additionally, it appears in the professional subplot, we see that the data also mostly conforms to a normal distribution, but with some more deviation in the upper tail. However, overall, these are small deviations so we shouldn't be too worried about violating normality assumptions.

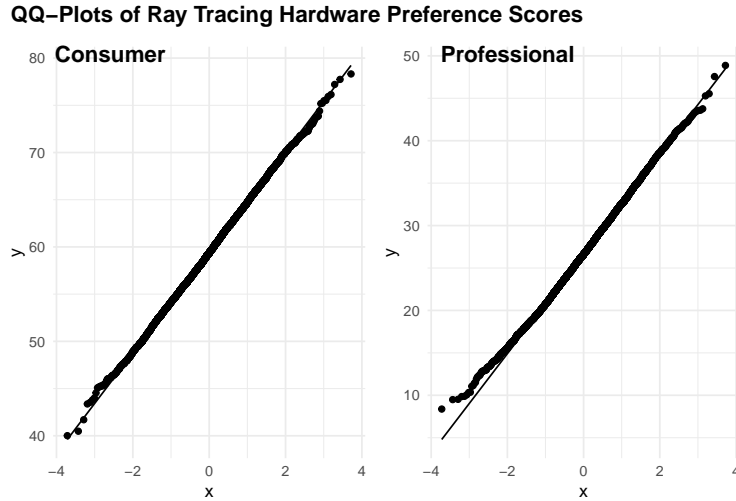


Figure 5: QQ-plots for Ray Tracing Hardware Preference Scores by User Type

The plot above depicts the QQ-plots for ray tracing hardware preference scores by user category. In the consumer subplot, we see that there are essentially no deviations from the QQ-plot line, i.e., we are not concerned about violating normality assumptions. In the professional subplot, we see that the only deviations from the QQ-plot line is in the lower tail, where the deviation is potentially indicative of a lighter than theoretical lower tail. Overall, we should not be very concerned about violating normality assumptions since any deviations from QQ-plot line are minor.

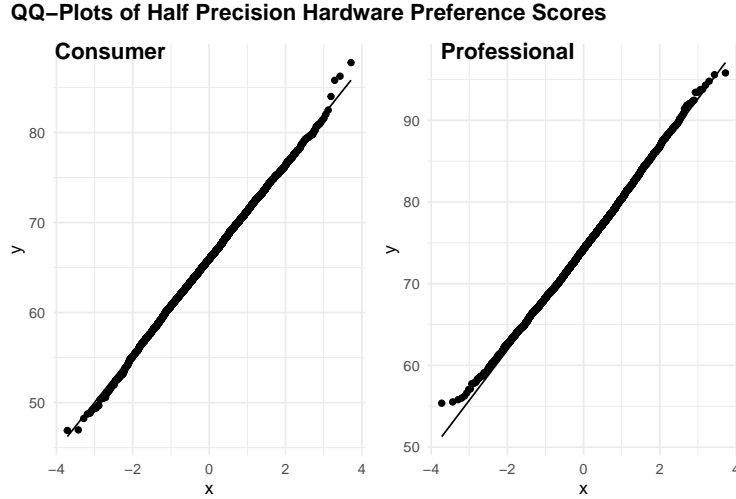


Figure 6: QQ-plots for Half Precision Hardware Preference Scores by User Type

The plot above depicts the QQ-plots for half precision matrix math hardware preference scores by user category. In the consumer subplot, we see that the only deviations from the QQ-plot line is in the upper tail, where the deviation is indicative of a slightly heavier than theoretical upper tail. In the professional subplot, we see a slight deviation from the QQ-plot line in the lower tail, where the small deviation is indicative of a slightly lighter than theoretical lower tail. Overall, we should not be very concerned about violating normality assumptions since any deviations from QQ-plot line are minor.

Then we examine the variances of the distributions by user type to see if we have evidence of unequal variances. For that we conduct 3 F-tests using the `var.test` command in R. We present a table with the preference score category, the test statistic, numerator degrees of freedom, denominator degrees of freedom, and the accompanying test p-value:

Table 5: F-test results for three hardware categories, comparing professional vs consumer users

Score Category	F-statistic	Numerator DF	Denominator DF	P-value
Traditional GPU	1.113758	5069	4929	0.0001411
Cores				
Ray-tracing	1.198882	5069	4929	$1.511 \cdot 10^{-10}$
Hardware				
Half Precision	1.304594	5069	4929	$< 2.2 \cdot 10^{-16}$
Hardware				

From the F-test results, we have extremely strong evidence in all three hardware categories of different variances between the distributions of professional and consumer users.

Since we feel comfortable with our normality assumption, i.e., our distributions are all approximately normal, we compute a confidence interval, using the R command `t.test`, for the differences in means between the three score categories separately.

Table 6: Confidence Interval for Traditional GPU Cores, Comparing Professional vs Consumer Users

Lower Bound of CI	Upper Bound of CI	Sample mean of Professional Users	Sample Mean of Consumer users
2.43740	2.89307	79.17382	76.50853

From our confidence interval examining the means of the traditional GPU cores preference scores of professionals and consumers, we find very strong evidence of evidence that the difference of means is not equal to 0. While the difference is small, i.e.,  $\approx 3$ , the p-value is less than  $2.2 \cdot 10^{-16}$ , i.e., very strong evidence of a difference in means. Thus, we conclude that it could be appropriate to use a mixed effects model for modelling the preference scores of traditional GPU cores.

Table 7: Confidence Interval for Ray Tracing Hardware, Comparing Professional vs Consumer Users

Lower Bound of CI	Upper Bound of CI	Sample mean of Professional Users	Sample Mean of Consumer users
-32.91641	-32.48230	26.70695	59.40630

Then from our confidence interval examining the difference in means of the ray tracing hardware preference scores of professionals and consumers, we a relatively small range from (-32.92, -32.48) for the true difference. The difference is very large, i.e.,  $\approx -33$ , and shows very strong statistically significant evidence that the mean preference score is different between professionals and consumers at the 95% significance level as 0 is not in the interval. Thus, we gather additional evidence that it is appropriate to use a mixed effects model for modelling the preference scores of ray tracing hardware.

Table 8: Confidence Interval for Half Precision Hardware, Comparing Professional vs Consumer Users

Lower Bound of CI	Upper Bound of CI	Sample mean of Professional Users	Sample Mean of Consumer users
8.143725	8.592351	74.32353	65.95549

Finally from our confidence interval of the difference in means of the half precision matrix math hardware preference scores of professionals and consumers, we find a relatively small range, i.e., 8.14 to 8.59. But the difference is also comparably large (compared to that of traditional GPU cores), i.e.,  $\approx 8$ , and shows very strong statistically significant evidence that the mean preference score is different between professionals and consumers at the 95% significance level as 0 is not in the interval. Thus, we conclude that it could also be appropriate to use a mixed effects model for modelling the preference scores of half precision matrix math hardware.

Since we are comfortable with our normality assumption and because the means of the distributions by user group are significantly different, meaning that we proceed with constructing a mixed effects model. The model architecture follows that as described in the methods section, and primarily our investigation here concerns if we should add a random slope to the weekly hours of GPU use variable for the user types. As mentioned in the methods section, we performed likelihood ratio tests for each of the three hardware categories, comparing the complex model with the extra random slope to the model without the random slope (reduced model). We find in all three cases that the complex model is not needed, i.e., the complex model does not provide a better fit.

For reference we fit the following models using the R package `lmer` [15] and all likelihood ratio tests were conducted using R package `lmtest` [16]:

- Traditional GPU Cores Preference Scores:
  - Reduced model: `model.tradGPUCores <- lmer(simulations.tradGPUCores ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType | user), data = sim.survey)`
  - Complex model: `model.extra.tradGPUCores <- lmer(simulations.tradGPUCores ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType + weeklyHours | user), data = sim.survey)`
- Ray Tracking Hardware Preference Scores:
  - Reduced model: `model.rayTraced <- lmer(simulations.rayTraced ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType | user), data = sim.survey)`
  - Complex model: `model.extra.rayTraced <- lmer(simulations.rayTraced ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType + weeklyHours | user), data = sim.survey)`
- Half Precision Matrix Math Hardware Preference Scores:
  - Reduced model: `model.halfPrecision <- lmer(simulations.halfPrecision ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType | user), data = sim.survey)`
  - Complex model: `model.extra.halfPrecision <- lmer(simulations.halfPrecision ~ weeklyHours + primaryInteraction + GPUPType + (1 + primaryInteraction + GPUPType + weeklyHours | user), data = sim.survey)`

Table 9: Traditional GPU Cores Models Likelihood Ratio Test

#Df	LogLik	Df	Chisq	Pr(>Chisq)
15	-30286.86	NA	NA	NA
11	-30285.96	-4	1.782022	0.7757699

The table above is the result of the likelihood ratio test between the complex model and the reduced model for the traditional GPU cores preference scores. As shown, the p-value is 0.775, so we have no evidence that the complex model for the traditional GPU cores preference scores being better. For conducting the likelihood ratio test, we use the following R function call `lmtest::lrtest(model.extra.tradGPUCores, model.tradGPUCores)`.

Table 10: Half Precision Hardware Models Likelihood Ratio Test

#Df	LogLik	Df	Chisq	Pr(>Chisq)
15	-30312.59	NA	NA	NA
11	-30309.00	-4	7.182054	0.1265746

The table above is the result of the likelihood ratio test between the complex model and the reduced model for the half precision hardware preference scores. As shown, the p-value is 0.126, so we have very little evidence that the complex model for the traditional GPU cores preference scores being better, which is insufficient to conclude that the complex model is necessary even at the 0.1 significance level. For conducting the likelihood ratio test, we use the following R function call `lmtest::lrtest(model.extra.halfPrecision, model.halfPrecision)`.

Table 11: Ray Tracing Hardware Models Likelihood Ratio Test

#Df	LogLik	Df	Chisq	Pr(>Chisq)
15	-30268.60	NA	NA	NA
11	-30268.83	-4	0.4514466	0.9780523

The table above is the result of the likelihood ratio test between the complex model and the reduced model for the ray tracing hardware preference scores. As shown, the p-value is 0.978, so we have essentially no evidence that the complex model for the traditional GPU cores preference scores being better. For conducting the likelihood ratio test, we use the following R function call `lmtest::lrtest(model.extra.rayTraced, model.rayTraced)`.

To summarize, our final models are the reduced complexity models as concluded from the likelihood ratio test. We present the following coefficients for the models. We will discuss the results comparatively for each model by user type because we want to examine the trends between GPU user groups. Note that for each of these tables, the coefficients (effect sizes) for the professional users are presented in the first row with consumers presented in the second row. First we will examine the coefficients for the traditional GPU cores preference score model coefficients:

	Intercept	Weekly Hours of GPU Usage	Primary GPU Interaction	GPU Type
0	60.30327	0.2471844	3.983912	1.865748
1	59.99129	0.2471844	3.876773	1.103765

From there, we see that it seems there is a difference between the coefficient estimate for the GPU Type between professionals and consumers. As in, GPU type contributes to traditional GPU cores preference scores almost twice as much in professionals (first row) when compared to consumers. In other words, professionals should expect to see around a 2 point increase for each GPU level they go up in (moving to more expensive/performant GPU categories), while consumers should only expect to see about a 1 point increase in their traditional GPU cores preference scores. All of the other effect sizes are extremely similar between professionals and consumers, e.g., the baseline preference score (intercept) for both is about 60. Additionally, our model estimates that for each hour extra of weekly GPU usage, both consumers and professionals will expect to see an  $\approx 0.25$  increase in preference for traditional GPU cores. Finally, for each increase in the “level” of primary GPU interaction (moving to less abstracted interactions with GPUs), both consumers and professionals should expect to see about a 4 point increase in their preferences for traditional GPU cores.

Then we examine the ray tracing hardware preference score model coefficients:

	Intercept	Weekly Hours of GPU Usage	Primary GPU Interaction	GPU Type
0	9.988032	0.2441839	4.110996	0.4393086
1	50.369219	0.2441839	1.922117	0.4716711

Immediately, we notice that the baseline preferences for ray tracing hardware is extremely different between professionals and consumers. Our models tells us to expect about a 10 for a baseline preference score for professionals while we should expect about a 50 for consumers. Another difference of note is that the coefficient estimate for primary GPU interaction is almost double in professionals when compared to consumers ( $\approx 4$  vs  $\approx 2$ ), i.e., professionals should expect to see about a 4 point increase in ray tracing hardware preference scores if they increase their primary GPU interaction to the next level (make their activities less abstracted away from the GPU), while consumers should only expect to see an increase of about 2. Finally both weekly hours of GPU usage and GPU type appear to be similar between professionals and consumers. Both professionals and consumers should expect to see an increase of  $\approx 0.24$  in their ray



tracing hardware scores for each additional hour of weekly GPU usage, while increasing the GPU quality they use by one level should expect to see an  $\approx 0.45$  increase to ray tracing hardware preference scores.

Finally we examine the half precision matrix math preference score model coefficients:

	Intercept	Weekly Hours of GPU Usage	Primary GPU Interaction	GPU Type
0	54.38693	0.2591262	4.139309	2.091543
1	54.38693	0.2591262	2.116416	1.037357

We notice that the baseline preference and the coefficient for weekly hours of GPU usage between professionals and consumers is the same for half precision hardware. So we expect that both user types have a baseline of about 54 points. While we also expect that for both user types, an hour increase in weekly GPU usage results in a  $\approx 0.26$  increase in half precision hardware preference score. However, we also see that the coefficient for primary GPU interaction for professionals is almost double that of consumers, i.e., for each increase in primary GPU interaction level (moving to activities less abstracted away from the GPU), we would expect to see an increase in 4 points for half precision hardware for professionals while consumers should only expect about 2. Then similarly, for GPU type, moving to the next higher quality GPU level, professionals should expect to see about a 2 point increase in half precision hardware scores while consumers should only expect to see about 1.

So overall, there are some interesting trends. Professionals and consumers appear to have about the same baseline preference for both traditional GPU cores and half precision hardware, but they greatly differ when it comes to ray tracing hardware with consumers greatly favouring ray tracing hardware relative to professionals. Generally, across hardware categories, the baseline preference for all hardware categories are fairly neutral (around 50), but it appears that professionals simply do not prefer ray tracing hardware as a baseline perhaps indicative of the lack of application in their domain [2]. Additionally, we see that in the cases of ray tracing hardware and half precision hardware, changing a professional’s primary GPU interaction will be much more impactful than doing the same with a consumer. This observation also occurs with GPU types for traditional GPU cores and half precision hardware. So it appears that professionals are more easily influenced into potentially liking a hardware category should they change up their hardware (a more powerful GPU) or their use cases when compared to a consumer, which may be indicative that professionals whose work requires a certain use case or quality of GPU prefer the hardware that their use case, GPU type combination entails. Finally, despite different opinions on different hardware categories, the hours a user uses their GPU seems to affect preference scores equally across hardware categories, suggesting that a particular hardware feature won’t affect the preferences of a user more strongly over other hardware features.

## Bibliography

1. Evanson, Nick. “Explainer: What Are Tensor Cores?” TechSpot, TechSpot, 27 July 2020, <https://www.techspot.com/article/2049-what-are-tensor-cores>.
2. Ravenscraft, Eric. “Should Anyone Actually Care about Ray Tracing?” Wired, Conde Nast, 3 Mar. 2021, <https://www.wired.com/story/should-anyone-actually-care-about-ray-tracing>.
3. Caulfield, Brian. “CPU vs GPU: What’s the Difference?” NVIDIA Blog, 23 June 2022, <https://blogs.nvidia.com/blog/2009/12/16/whats-the-difference-between-a-cpu-and-a-gpu>.
4. “GPU Processing Unit (GPU) Market” Allied Market Research, <https://bit.ly/3BLAyWQ>.
5. “AMD Reports First Quarter 2022 Financial Results.” Advanced Micro Devices, Inc., 3 May 2022, <https://ir.amd.com/news-events/press-releases/detail/1062/amd-reports-first-quarter-2022-financial-results>.
6. “Nvidia Announces Financial Results for First Quarter Fiscal 2023.” NVIDIA Newsroom, 20 Sept. 2022, <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-first-quarter-fiscal-2023>.
7. Pandey, Mohit, et al. “The Transformational Role of GPU Computing and Deep Learning in Drug Discovery.” Nature News, Nature Publishing Group, 23 Mar. 2022, <https://www.nature.com/articles/s42256-022-00463-x>.

8. CSLab Support, <https://support.cs.toronto.edu/systems/linuxgpu.html>.
9. “Games Industry Data and Analysis.” Video Game Insights, <https://vginsights.com/insights/article/report-steam-games-market-size-likely-to-decline-in-2022-after-reaching-6-6bn-in-2021>.
10. Combs, Veronica, et al. “8 Hours and 27 Minutes. That’s How Long the Average Gamer Plays Each Week.” TechRepublic, 22 Sept. 2022, <https://www.techrepublic.com/article/8-hours-and-27-minutes-thats-how-long-the-average-gamer-plays-each-week/>.
11. Ghasemi, Asghar, and Saleh Zahediasl. “Normality tests for statistical analysis: a guide for non-statisticians.” *International journal of endocrinology and metabolism* vol. 10,2 (2012): 486-9. doi: 10.5812/ijem.3505
12. “Best Buy Midrange GPU.” BestBuy.com, <https://www.bestbuy.com/site/shop/best-mid-range-graphics-card>.
13. “Best Buy High End GPU.” BestBuy.com, <https://www.bestbuy.com/site/shop/high-end-graphics-cards>.
14. “Difference in Means.” Calcworkshop, 10 Oct. 2020, <https://calcworkshop.com/confidence-interval/difference-in-means/>.
15. Bates, D., M. Mächler, B. Bolker, and S. Walker. “Fitting Linear Mixed-Effects Models Using Lme4”. *Journal of Statistical Software*, vol. 67, no. 1, Oct. 2015, pp. 1-48, doi:10.18637/jss.v067.i01.
16. “Testing Linear Regression Models [R Package Lmtest Version 0.9-40].” The Comprehensive R Archive Network, Comprehensive R Archive Network (CRAN), 21 Mar. 2022, <https://cran.r-project.org/web/packages/lmtest/index.html>.
17. Xie, Yihui. “A General-Purpose Package for Dynamic Report Generation in R [R Package Knitr Version 1.40].” The Comprehensive R Archive Network, Comprehensive R Archive Network (CRAN), 24 Aug. 2022, <https://cran.r-project.org/web/packages/knitr/index.html>.
18. Martindale, Jon, et al. “Nvidia RTX DLSS: Everything You Need to Know.” Digital Trends, Digital Trends, 20 Sept. 2022, <https://www.digitaltrends.com/computing/everything-you-need-to-know-about-nvidias-rtx-dlss-technology/>.
19. “Automatic Mixed Precision for Deep Learning.” NVIDIA Developer, 18 Feb. 2022, <https://developer.nvidia.com/automatic-mixed-precision>.
20. “Package Extradistr.” CRAN, Comprehensive R Archive Network (CRAN), <https://cran.r-project.org/web/packages/extraDistr/index.html>.

## Appendix

Here is a glimpse of the data set simulated:

```
## Rows: 10,000
## Columns: 10
## $ user <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,~
## $ primaryInteraction <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 3, 3, 3, 3, 3,~
## $ GPUType <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 4, 4, 4, 4, 4,~
## $ weeklyHours <dbl> 35.895067, 35.895067, 35.895067, 35.895067, ~
## $ ce.tradGPUCores <dbl> 74.97377, 74.97377, 74.97377, 74.97377, 74.9~
## $ ce.rayTraced <dbl> 23.47377, 23.47377, 23.47377, 23.47377, 23.4~
## $ ce.halfPrecision <dbl> 69.97377, 69.97377, 69.97377, 69.97377, 69.9~
## $ simulations.tradGPUCores <dbl> 71.57973, 77.84533, 71.45119, 72.30385, 78.8~
## $ simulations.rayTraced <dbl> 33.68354, 26.54617, 25.58342, 20.99166, 25.9~
## $ simulations.halfPrecision <dbl> 76.18876, 79.88482, 66.74156, 74.80471, 62.8~
```